

GALILEO protocol (ground-truth multi-turn robustness)

Phase 1: Initial evaluation

Answer once on ground-truth tasks
Score correctness → keep initially-correct set
Export per-example logs

logs + C

Phase 2: Multi-round pressure

For each $x \in C$, run up to R rounds
Persona pressure (5 personas) OR
Neutral Re-asking Control (drift baseline)
Measure survival + TOF per round

flip set

Phase 3: Recovery

For flipped examples, apply recovery prompt
Score recovery conditional on flip
Report recovery as distinct axis

Outputs (per run)

paper_exports/: survival_curve.csv • turn_of_failure.csv • flip_samples.csv • metadata.json • runner_metadata.json

Tracked artifacts: compact CSV summaries under docs/paper/artifacts/ → submission figures under docs/paper/figures/